

# Generative Adversarial Networks

CMU 11-785 Introduction to Deep Learning, Fall 2020  
lecture 24

Presenter : Changdae Oh  
bnormal16@naver.com

TAVE Research DL001  
2021.05.30

# Topics

1. Recap
2. GAN Optimization Issues
3. GAN Training & Stabilization
4. Conclusion

# Topics

1. Recap
2. GAN Optimization Issues
3. GAN Training & Stabilization
4. Conclusion

## Adversarial Learning

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

$$L = -(t \log(y) + (1 - t) \log(1 - y))$$

**Discriminator** calculates a divergence between generated and target distribution.  
acts as a loss function.

**Generator** tries to minimize the divergence.

- Powerful tool for generative modeling
- Lots of potential
- Limited by pragmatic issues (stability)

# Recap

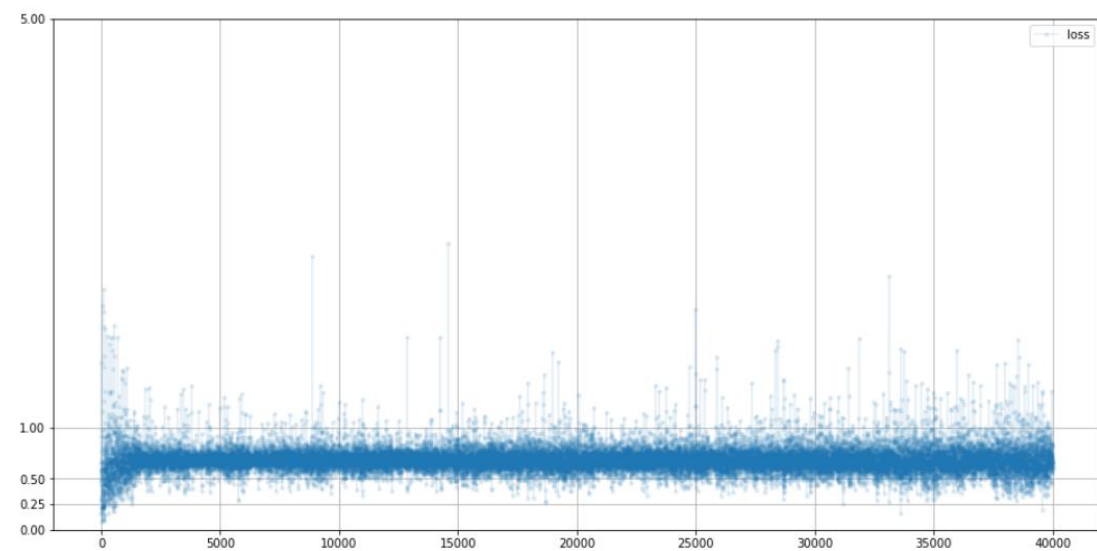
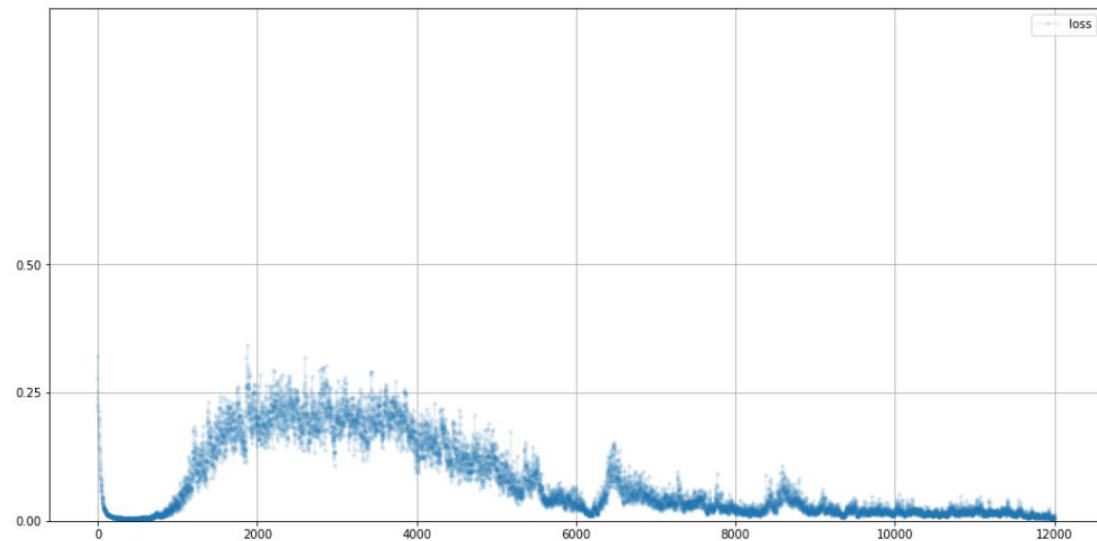
```
1 class Discriminator(nn.Module):
2
3     def __init__(self):
4         super().__init__()
5         self.model = nn.Sequential(
6             nn.Conv2d(3, 256, kernel_size=8, stride=2),
7             nn.BatchNorm2d(256),
8             nn.GELU(),
9
10            nn.Conv2d(256, 256, kernel_size=8, stride=2),
11            nn.BatchNorm2d(256),
12            nn.GELU(),
13
14            nn.Conv2d(256, 3, kernel_size=8, stride=2),
15            nn.GELU(),
16
17            View(3*10*10),
18            nn.Linear(3*10*10, 1),
19            nn.Sigmoid()
20        )
21
22        self.loss_function = nn.BCELoss()
23        self.optimizer = torch.optim.Adam(self.parameters(), lr=0.0001)
24
25
26 def train(self, inputs, targets):
27     outputs = self.forward(inputs)
28     loss = self.loss_function(outputs, targets)
29
30     self.optimizer.zero_grad()
31     loss.backward()
32     self.optimizer.step()
```

```
1 class Generator(nn.Module):
2     def __init__(self):
3         super().__init__()
4
5         self.model = nn.Sequential(
6             nn.Linear(100, 3*11*11),
7             nn.GELU(),
8             View((1, 3, 11, 11)),
9
10            nn.ConvTranspose2d(3, 256, kernel_size=8, stride=2),
11            nn.BatchNorm2d(256),
12            nn.GELU(),
13
14            nn.ConvTranspose2d(256, 256, kernel_size=8, stride=2),
15            nn.BatchNorm2d(256),
16            nn.GELU(),
17
18            nn.ConvTranspose2d(256, 3, kernel_size=8, stride=2, padding=1),
19            nn.BatchNorm2d(3),
20            nn.Sigmoid()
21        )
22
23        self.optimizer = torch.optim.Adam(self.parameters(), lr=0.0001)
24
25
26 def train(self, D, inputs, targets):
27     g_output = self.forward(inputs)
28     d_output = D.forward(g_output)
29     loss = D.loss_function(d_output, targets)
30
31     self.optimizer.zero_grad()
32     loss.backward()
33     self.optimizer.step()
```

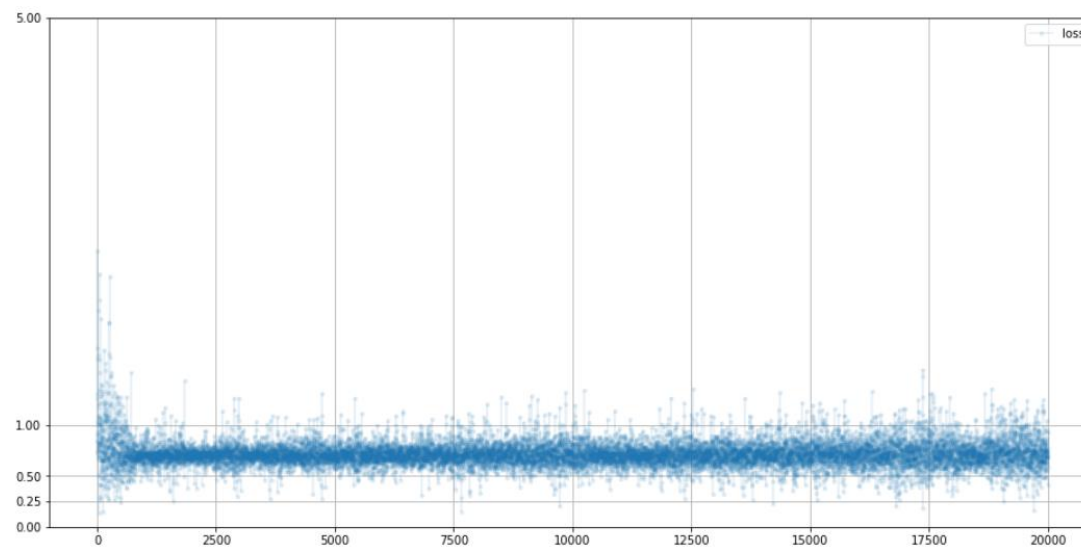
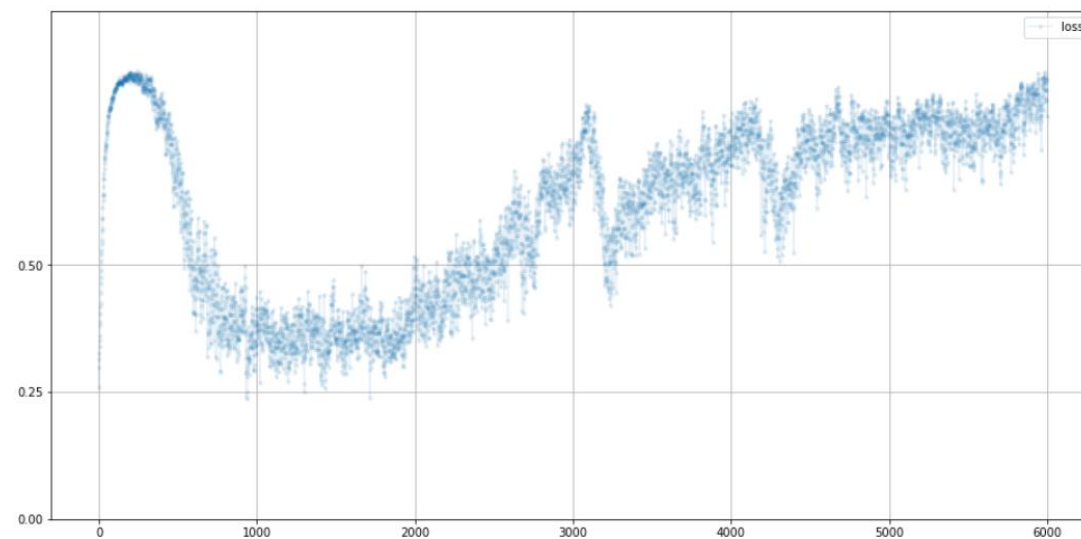
## Train Loop

```
1 %%time
2 epochs = 10
3 D = Discriminator() ; D.to(device)
4 G = Generator() ; G.to(device)
5
6 for epoch in range(epochs):
7     print ("epoch = ", epoch + 1)
8
9     for image_data_tensor in celeba_dataset:
10         D.train(image_data_tensor, torch.cuda.FloatTensor([1.0]))
11         D.train(G.forward(generate_random_seed(100)).detach(), torch.cuda.FloatTensor([0.0]))
12         G.train(D, generate_random_seed(100), torch.cuda.FloatTensor([1.0]))
```

## Discriminator



## Generator

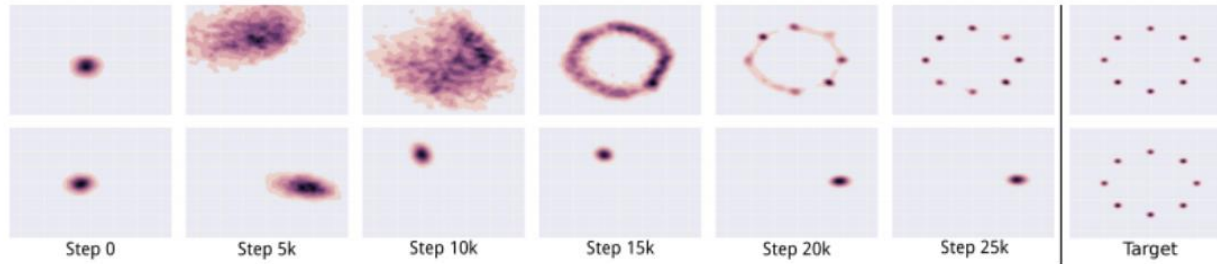


# Topics

1. Recap
- 2. GAN Optimization Issues**
3. GAN Training & Stabilization
4. Conclusion

# GAN Optimization Issues

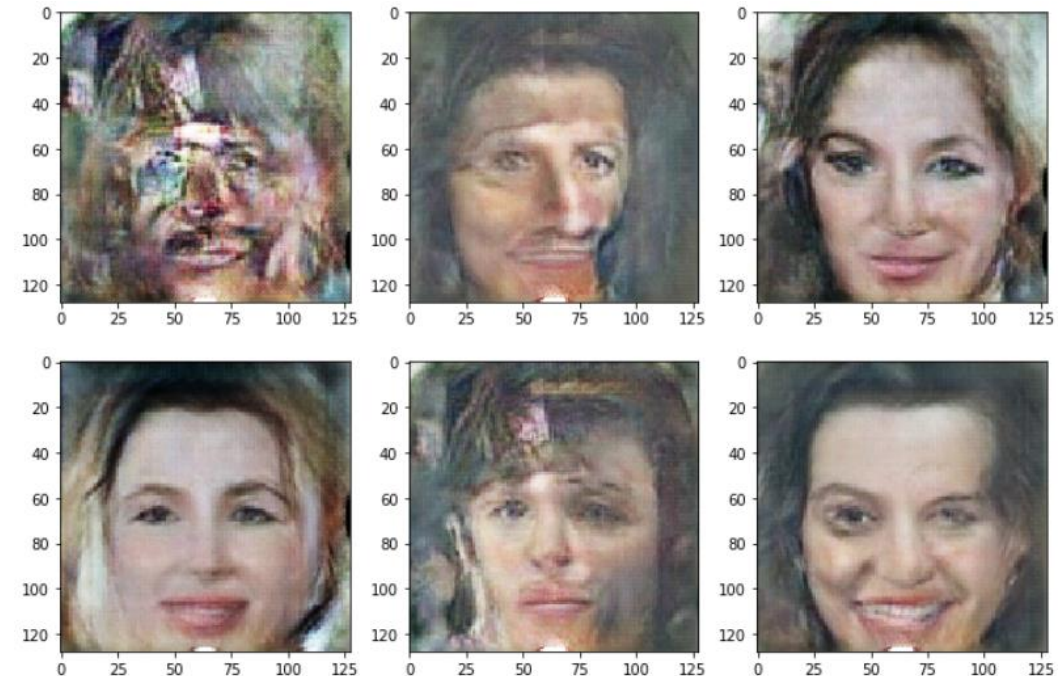
## Mode Collapse



- Cause can be unclear
  - Unbalanced complexity?
  - minimax  $\leftrightarrow$  maximin
  - Discriminator saturation
  - Information about image diversity, quality is not reflected in the loss function.



worse



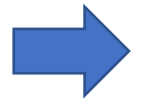
better



# GAN Optimization Issues

How about just train an optimal discriminator/generator?

- The optimal discriminator emits 0.5 for all inputs...
- Optimal discriminator conditional on current generator and vice-versa.
- Cannot train generator without training discriminator first



G. and D. must be trained together!

- Simultaneous updates require a careful **balance** between players.
- There is a stationary point but no guarantee of reaching it.

## Factors Affecting Balance

- Different optimizers and learning rates
- Different architectures, depths, # of parameters
- Regularization
- Iterations

## Convergence in two-player games

- Rock, Paper, Scissors
  - Global optimum - Both :  $(0.33, 0.33, 0.33)$ 
    - Both player win, lose or draw w.p. 0.33
  - Local optimum - D :  $(0.4, 0.3, 0.3)$  , G :  $(0, 1, 0)$ 
    - G wins, loses or draws w.p.  $(0.4, 0.3, 0.3)$

# Topics

1. Recap
2. GAN Optimization Issues
- 3. GAN Training & Stabilization**
4. Conclusion

## GAN Training Techniques

### Adjusted / Regularized gradient

- UGAN
- Gradient descent is locally stable
- DRAGAN
- Numerics of GANs

### Modified loss function

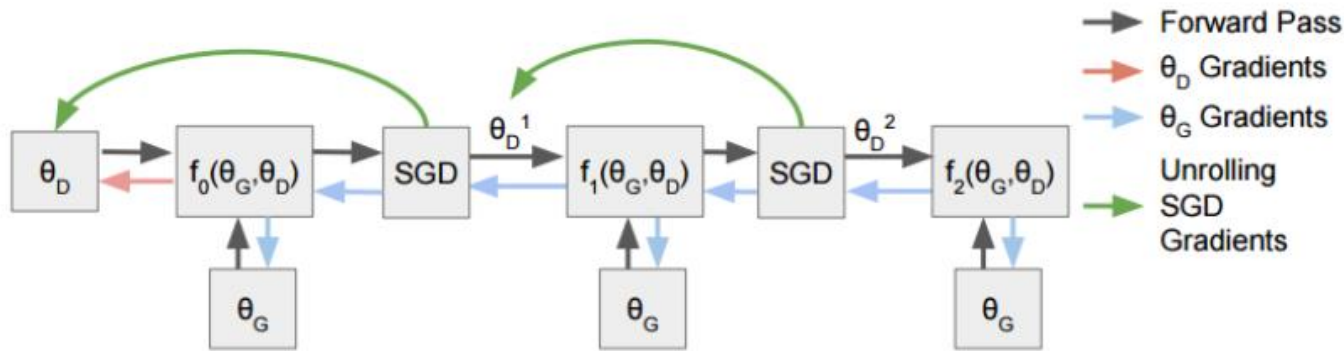
- LSGAN
- EBGAN
- WGAN
- WGAN – GP
- SNGAN

### Others

- Instance Noise
- ProGAN
- Other improved techniques for GANs

*It is a field that is currently being studied very actively,  
and the dynamics of GANs have not been fully understood !*

- Look ahead updating : optimize future loss, not current loss



$$f(\theta_G, \theta_D) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x; \theta_D)] + \mathbb{E}_{z \sim p_x(z)} [\log(1 - D(G(z; \theta_G); \theta_D))]$$

$$\theta_D^0 = \theta_D$$

$$\theta_D^{k+1} = \theta_D^k + \eta^k \frac{df(\theta_G, \theta_D^k)}{d\theta_D^k}$$

$$f_K(\theta_G, \theta_D) = f(\theta_G, \theta_D^K(\theta_G, \theta_D))$$

Actual update :

$$\theta_G \leftarrow \theta_G - \eta \frac{df_K(\theta_G, \theta_D)}{d\theta_G}$$

$$\theta_D \leftarrow \theta_D + \eta \frac{df(\theta_G, \theta_D)}{d\theta_D}.$$

### GD is locally stable

$$\ell_G = \ell_{G,0} + \eta \|\nabla \ell_D\|^2$$

Minimize the original objective as well as the gradient of the discriminator.

### DRAGAN

$$\lambda \mathbb{E}_{X, \epsilon} \max(0, \|\nabla D(X + \epsilon)\|^2 - k)$$

Minimize the norm of the gradient in a region around real data. (gradient penalty)

### The numerics of GANs

$$L = L_0 + \lambda \|\nabla L_0\|_2^2$$

G & D Mutual de-escalation

# Improved Techniques

## Feature matching

$$\|E_{x \sim p_{data}} f(x) - E_{z \sim p_z(z)} f(G(z))\|_2^2$$

Statistics of generated images should match statistics of real images

## Historical averaging

$$\left\| \theta - \frac{1}{t} \sum_i^t \theta_i \right\|_2^2$$

Dampen oscillations by encouraging Updates to converge to a mean

# GAN Training & Stabilization

## Minibatch discrimination

Discriminator can look at multiple inputs at once and decide if those inputs come from the real or generated distribution

➡ Make robust model

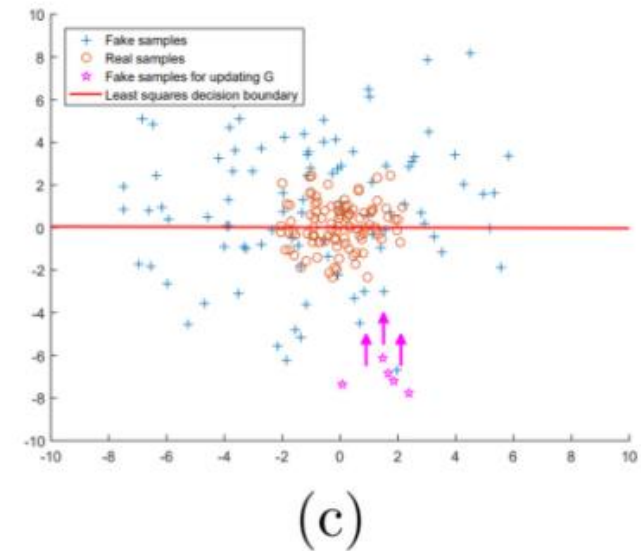
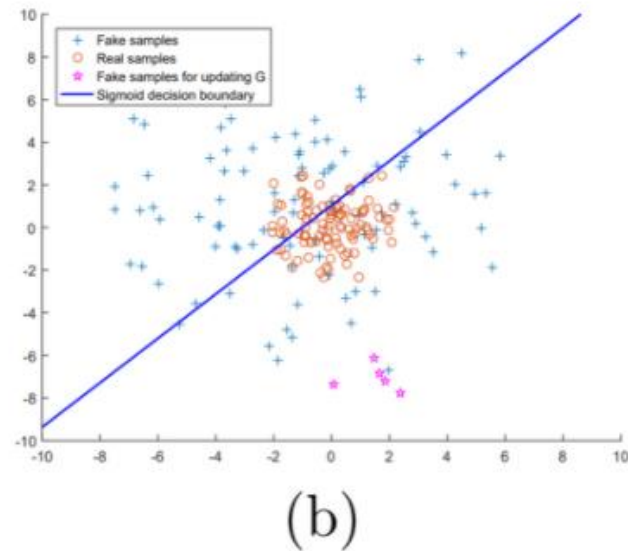
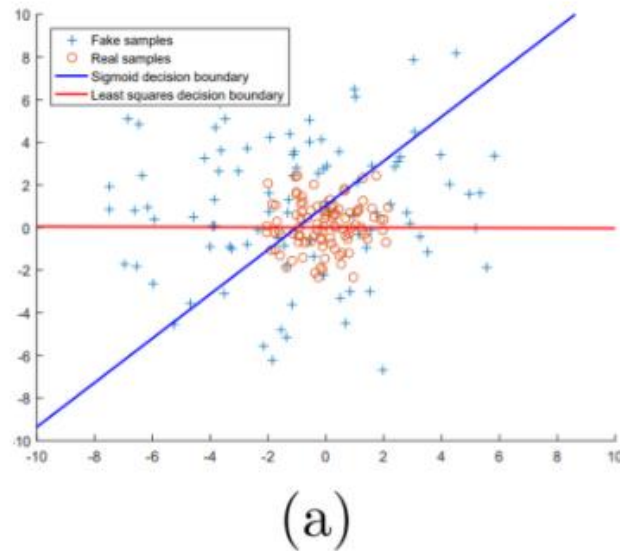
## Label smoothing

- Common and easy technique that improves performance across many domains
- Smooth the real targets but do not smooth the generated targets when training the discriminator

## Least-Squares GAN

$$\min_D V_{\text{LSGAN}}(D) = \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log(D(x) - b)^2] + \frac{1}{2} \mathbb{E}_{z \sim p_z(z)} [\log(D(G(z)) - a)^2]$$
$$\min_G V_{\text{LSGAN}}(G) = \frac{1}{2} \mathbb{E}_{z \sim p_z(z)} [\log(D(G(z)) - c)^2]$$

Use an L2 loss instead of cross-entropy





# Modified loss function

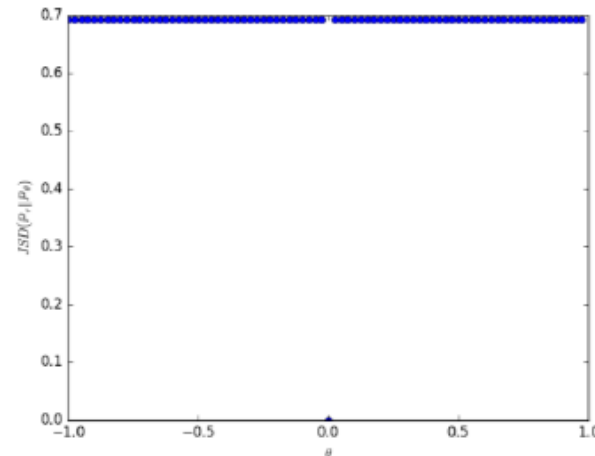
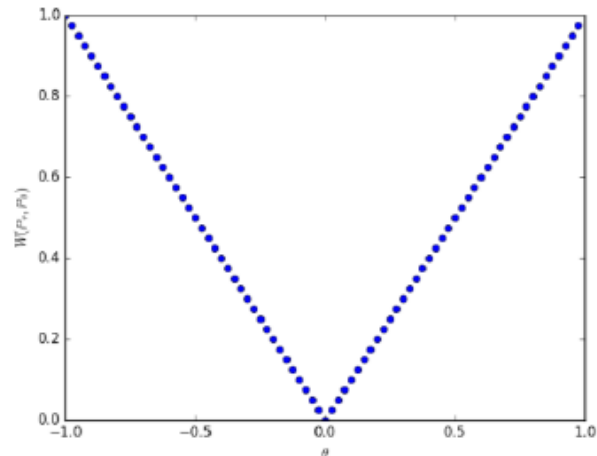
# GAN Training & Stabilization

## WGAN

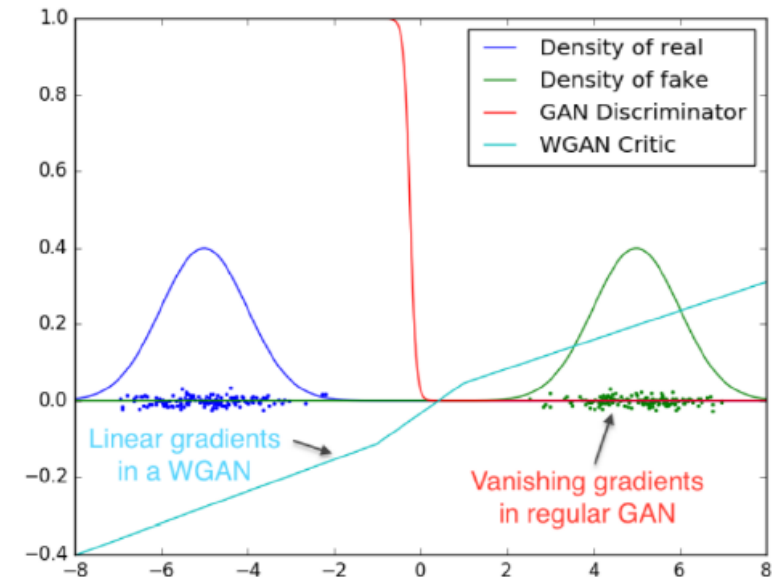
$$L_D = \mathbb{E}_X D(X) - \mathbb{E}_Z D(G(Z))$$

$$L_G = \mathbb{E}_Z D(G(Z))$$

- Simplify the loss function
  - Eliminate log term
  - Wasserstein Loss
- Drop sigmoid term from discriminator
- D must be 1-Lipschitz continuous function
  - Abs of the slope must be up to 1 anywhere.



- $W(P_0, P_\theta) = |\theta|$ ,
- $JS(P_0, P_\theta) = \begin{cases} \log 2 & \text{if } \theta \neq 0, \\ 0 & \text{if } \theta = 0, \end{cases}$
- $KL(P_\theta || P_0) = KL(P_0 || P_\theta) = \begin{cases} +\infty & \text{if } \theta \neq 0, \\ 0 & \text{if } \theta = 0, \end{cases}$
- and  $\delta(P_0, P_\theta) = \begin{cases} 1 & \text{if } \theta \neq 0, \\ 0 & \text{if } \theta = 0. \end{cases}$



## WGAN (compare to GAN)

- Use Wasserstein loss
- True : 1, Fake : -1
- Not use sigmoid on top of discriminator
- Clipping the weight of the discriminator after each update  
(To satisfy the Lipschitz constraint.)
- Each time we update the generator, we train the discriminator several times.

Get a stronger discriminator

And can better balance the training of discriminators and constructors.

However,

Training speed is significantly reduced because weights are clipped in the discriminator.

## WGAN - GP

$$L = \mathbb{E}_X(D(X)) - \mathbb{E}_Z(D(G(Z))) + \lambda \mathbb{E}_{X'} (\|\nabla D(X')\|_2 - 1)^2$$

- Demonstrate the flaws with gradient clipping
- Calculate the gradient of D at random samples
- Use samples that are random linear interpolations between real and fake
- Add a penalty of the mean squared distance between the gradient and 1

---

**WGAN** : bound-based – can easily learn poor local optima

**WGAN GP** : sampling based – can be unreliable

# Topics

1. Recap
2. GAN Optimization Issues
3. GAN Training & Stabilization
- 4. Conclusion**

## Current & Trend

- Regularize & smooth the discriminator
  - Update the players towards some sort of consensus or future stable region, not greedy
  - Simplify networks and losses to eliminate nonlinearities
  - Constrain the complexity of the discriminator
- 
- Lots of GPUs and tuning required
  - better techniques allow larger models, always pushing the limits
  - Mostly images but other modalities are in-progress

- How can we measure / constrain the complexity of a NN?
- Evaluation
- How can we architect a neural network to learn a meaningful loss?  
(Perceptual distance between things )

## “The GAN Zoo”

- GAN - Generative Adversarial Networks
- 3D-GAN - Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling
- acGAN - Face Aging With Conditional Generative Adversarial Networks
- AC-GAN - Conditional Image Synthesis With Auxiliary Classifier GANs
- AdaGAN - AdaGAN: Boosting Generative Models
- AEGAN - Learning Inverse Mapping by Autoencoder based Generative Adversarial Nets
- AffGAN - Amortised MAP Inference for Image Super-resolution
- AL-CGAN - Learning to Generate Images of Outdoor Scenes from Attributes and Semantic Layouts
- ALI - Adversarially Learned Inference
- AM-GAN - Generative Adversarial Nets with Labeled Data by Activation Maximization
- AnoGAN - Unsupervised Anomaly Detection with Generative Adversarial Networks to Guide Marker Discovery
- ArtGAN - ArtGAN: Artwork Synthesis with Conditional Categorical GANs
- b-GAN - b-GAN: Unified Framework of Generative Adversarial Networks
- Bayesian GAN - Deep and Hierarchical Implicit Models
- BEGAN - BEGAN: Boundary Equilibrium Generative Adversarial Networks
- BiGAN - Adversarial Feature Learning
- BS-GAN - Boundary-Seeking Generative Adversarial Networks
- CGAN - Conditional Generative Adversarial Nets
- CaloGAN - CaloGAN: Simulating 3D High Energy Particle Showers in Multi-Layer Electromagnetic Calorimeters with Generative Adversarial Networks
- CCGAN - Semi-Supervised Learning with Context-Conditional Generative Adversarial Networks
- CatGAN - Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks
- CoGAN - Coupled Generative Adversarial Networks
- Context-RNN-GAN - Contextual RNN-GANs for Abstract Reasoning Diagram Generation
- C-RNN-GAN - C-RNN-GAN: Continuous recurrent neural networks with adversarial training
- CS-GAN - Improving Neural Machine Translation with Conditional Sequence Generative Adversarial Nets
- CVAE-GAN - CVAE-GAN: Fine-Grained Image Generation through Asymmetric Training
- CycleGAN - Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks
- DTN - Unsupervised Cross-Domain Image Generation
- DCGAN - Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks
- DiscoGAN - Learning to Discover Cross-Domain Relations with Generative Adversarial Networks
- DR-GAN - Disentangled Representation Learning GAN for Pose-Invariant Face Recognition
- DualGAN - DualGAN: Unsupervised Dual Learning for Image-to-Image Translation
- EBGAN - Energy-based Generative Adversarial Network
- f-GAN - f-GAN: Training Generative Neural Samplers using Variational Divergence Minimization
- FF-GAN - Towards Large-Pose Face Frontalization in the Wild
- GAWWN - Learning What and Where to Draw
- GeneGAN - GeneGAN: Learning Object Transfiguration and Attribute Subspace from Unpaired Data
- Geometric GAN - Geometric GAN
- GoGAN - Gang of GANs: Generative Adversarial Networks with Maximum Margin Ranking
- GP-GAN - GP-GAN: Towards Realistic High-Resolution Image Blending
- IAN - Neural Photo Editing with Introspective Adversarial Networks
- iGAN - Generative Visual Manipulation on the Natural Image Manifold
- IcGAN - Invertible Conditional GANs for image editing
- ID-CGAN - Image De-raining Using a Conditional Generative Adversarial Network
- Improved GAN - Improved Techniques for Training GANs
- InfoGAN - InfoGAN: Interpretable Representation Learning by Information Maximizing Generative Adversarial Nets
- LAGAN - Learning Particle Physics by Example: Location-Aware Generative Adversarial Networks for Physics Synthesis
- LAPGAN - Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks